

Embedded Video Systems With Zephyr

Josuah Demangeon, Panoramix Labs, tinyVision.ai

Background

A contractor working essentially with tinyVision.ai



Autoportrait: I am curious about a lot of topics, but only scratch the surface.

You are welcome and invited to dive in depth!

Video Systems

Famous example: home cinema

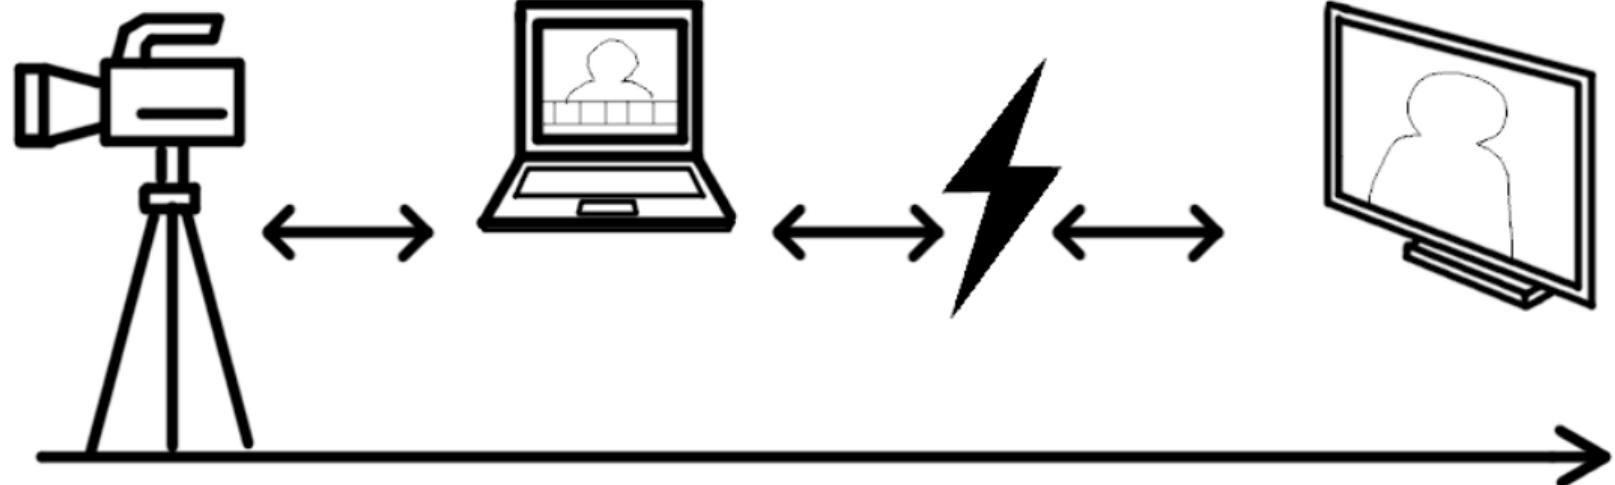


Pro Video
Camera

Video editing
software

Distribution
network

Final destination:
human eyes

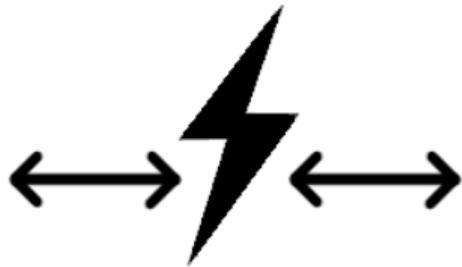


Pro Video
Camera

Video editing
software

Distribution
network

Final destination:
human eyes

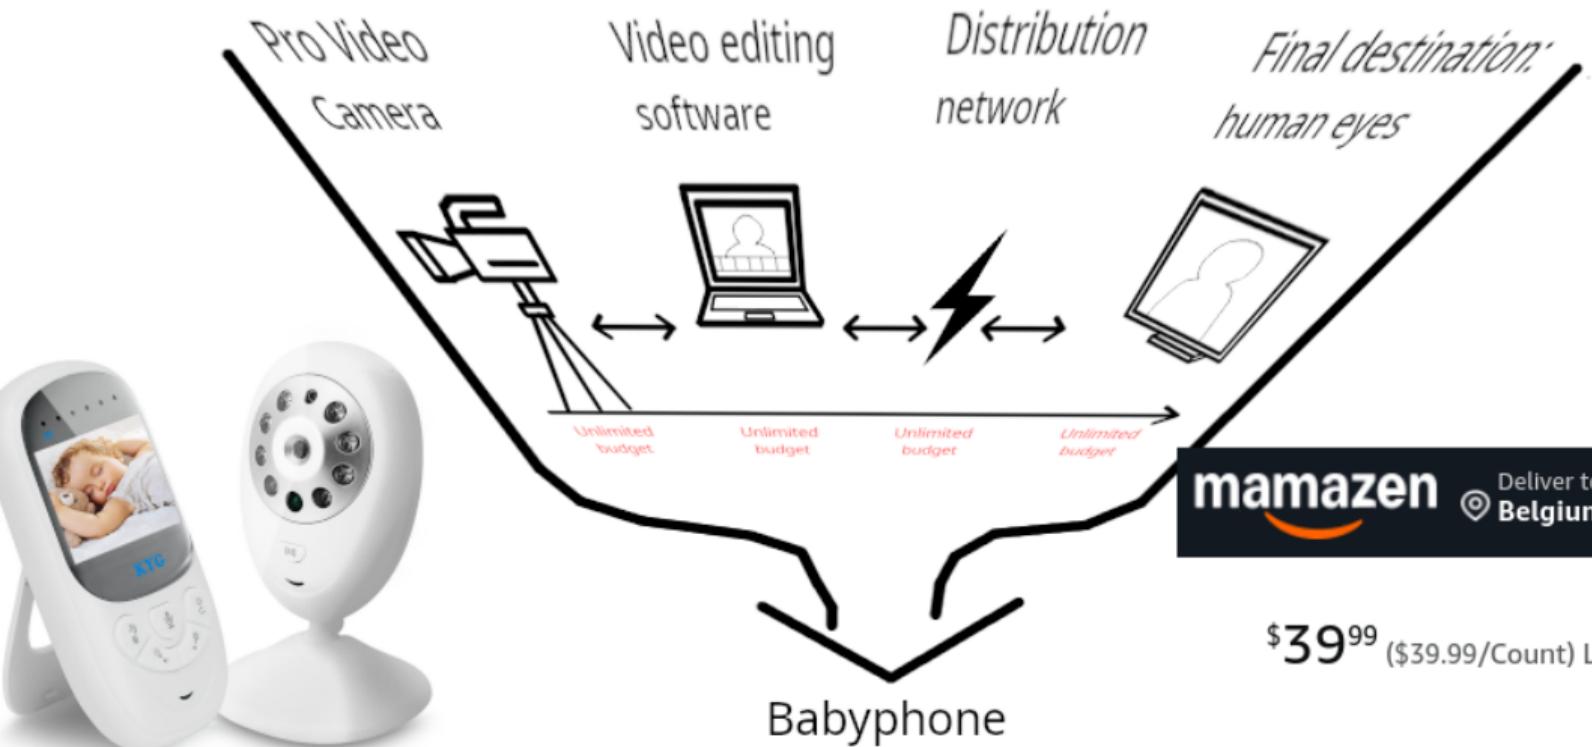


Unlimited
budget

Unlimited
budget

Unlimited
budget

Unlimited
budget



mamazen Deliver to Belgium

\$39⁹⁹ (\$39.99/Count) List: \$49.99

All ba

Embedded Video Systems

Constraints:

- Cost budget
- Processing budget
- Time budget (low-latency, real-time)

Can only work at low-resolution...

Embedded Video Systems

Constraints:

- Cost budget
- Processing budget
- Time budget (low-latency, real-time)

Can only work at low-resolution... <- FALSE!

Pro Video
Camera



Video editing
software



Distribution
network



Final destination:
human eyes



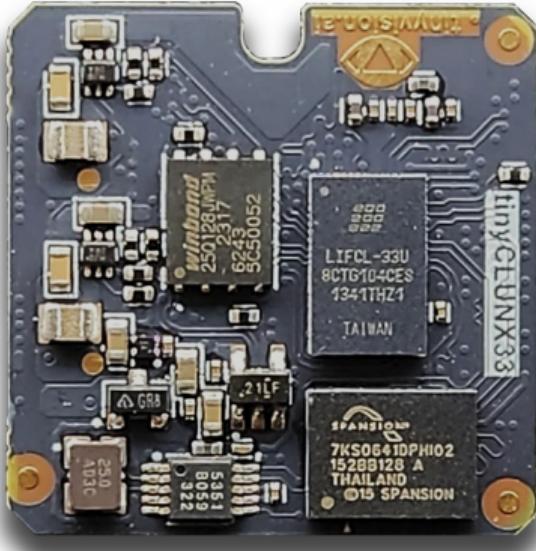
Needs an operating system too!

Embedded is not always low-end.

Embedded Video Systems

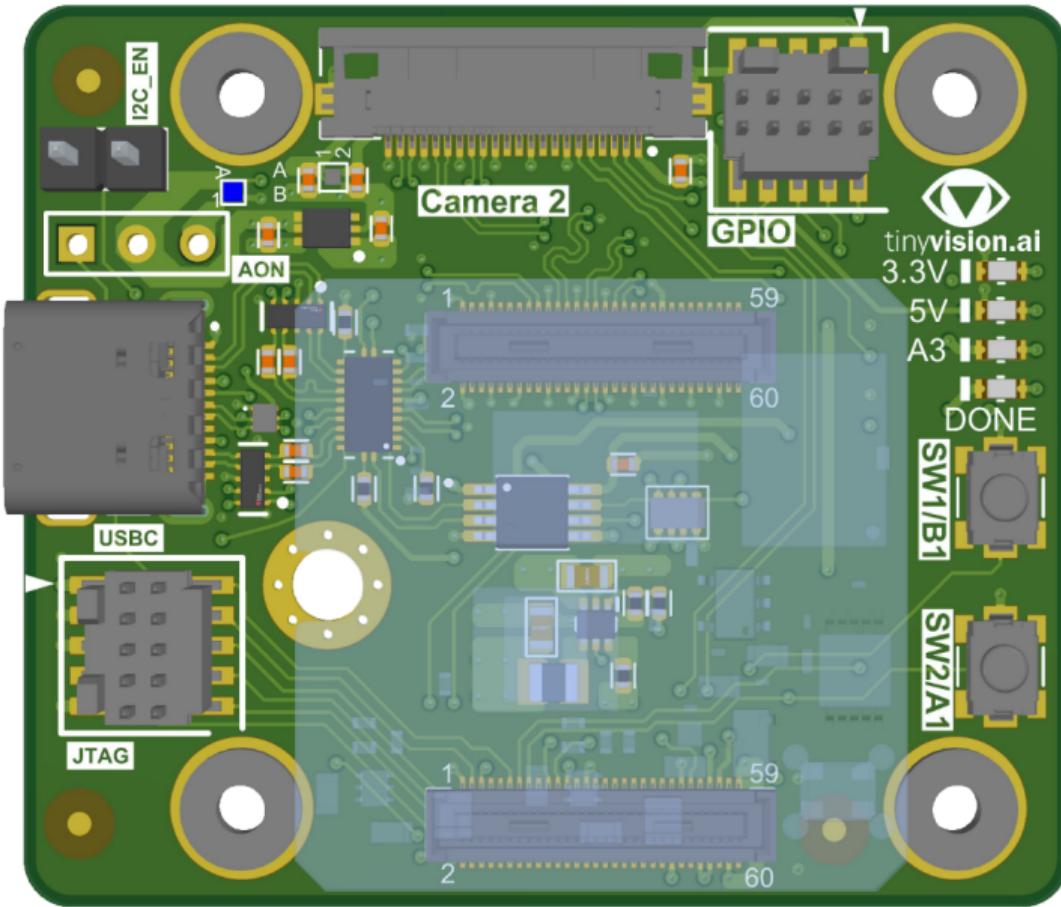
"Why not use an USB camera?"

We are now implementing the USB camera *itself*.



tinyCLUNX33: the heart of an USB 3 webcam.

3.4 Gbit/s under 80 mW



Embedded Video Systems

"Why not just a Raspberry Pi?"

→ Power budget

→ Performance

→ Cost

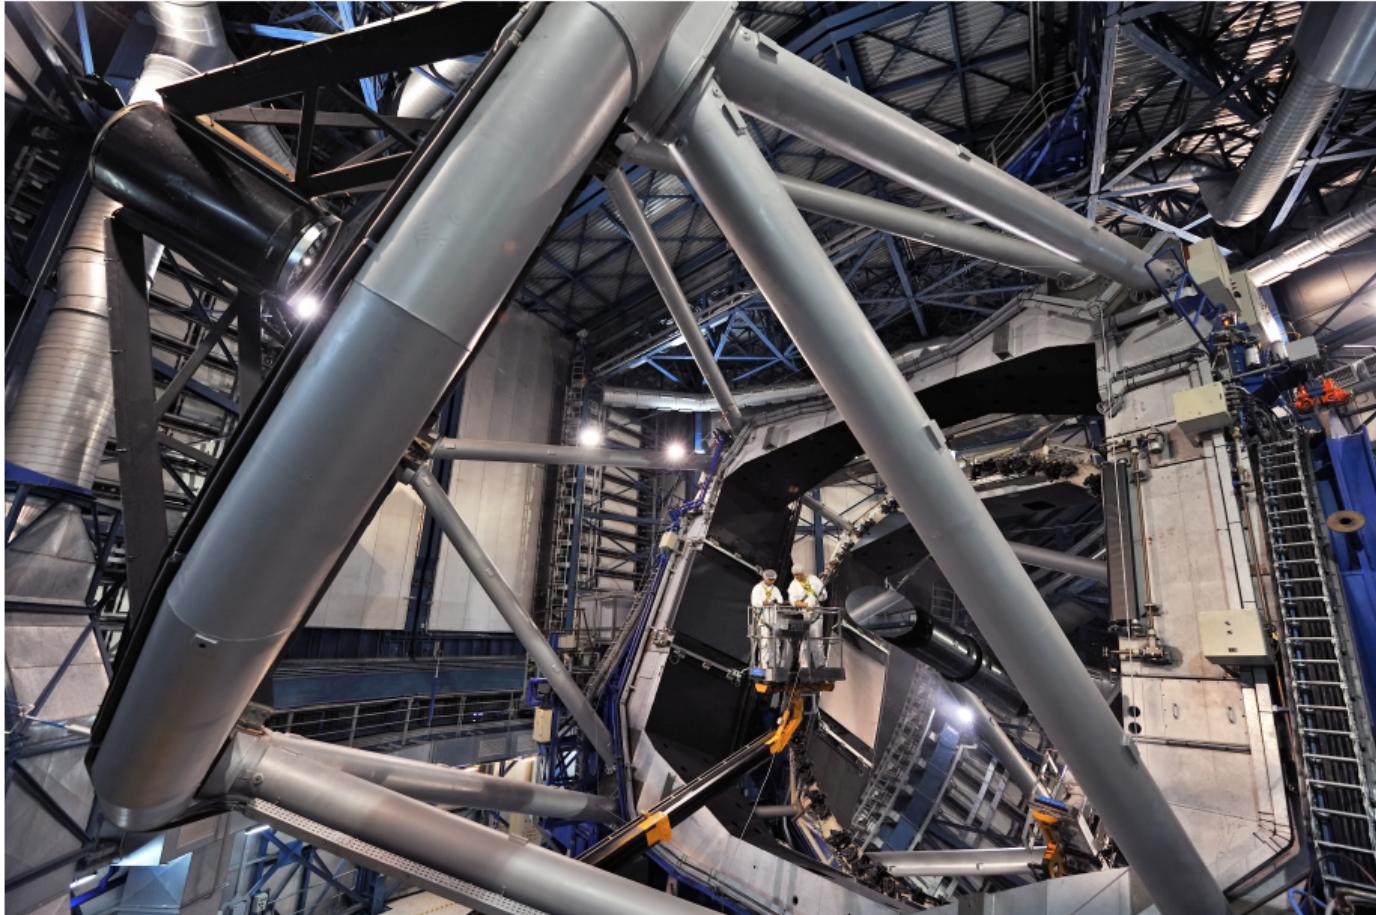
→ Latency



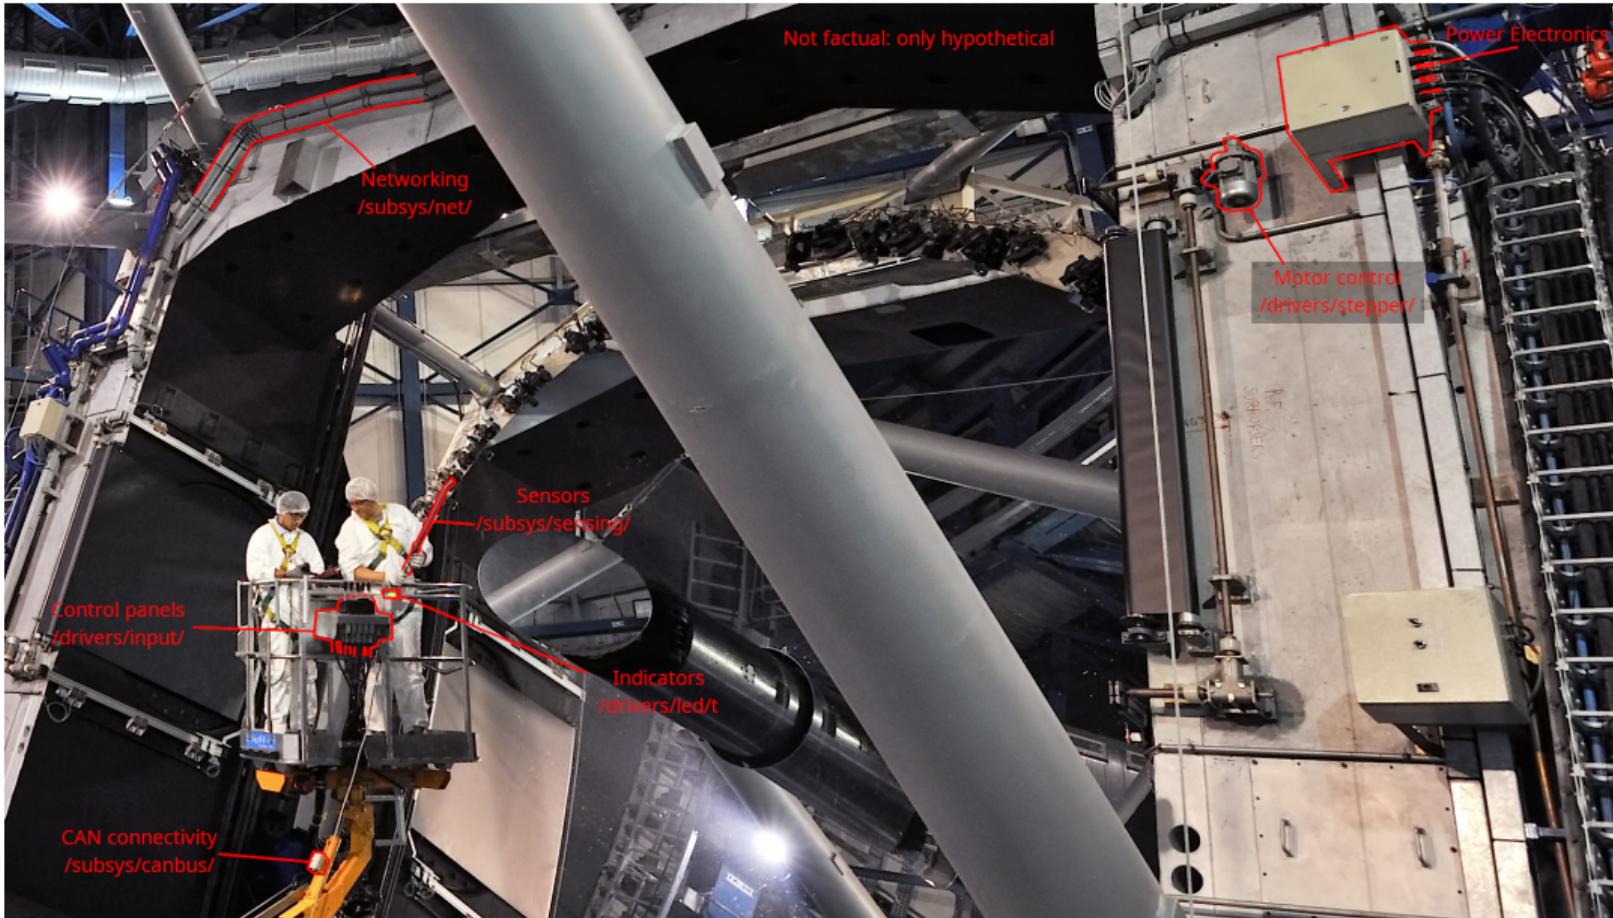
Embedded Video Systems

Can be very large:





We can imagine a lot involved to assist the video function:



Still there on small embedded systems:

→ Motor for auto-focus ("VCM" motor

```
#include <zephyr/drivers/video-controls.h>)
```

→ I2C communication with other chips

```
(#include <zephyr/drivers/i2c.h>)
```

→ Turning on/off the chip power ([Power Management](#))

Embedded Video Systems

But usually the smaller the better: how to shrink?

Switch from Linux OS → RTOS like Zephyr

Which board running it can do video?

Video-capable boards running Zephyr

Video-capable means DVP (a.k.a. parallel port) or MIPI support.

Different software ecosystem

FFmpeg → ???

Gstreamer → ???

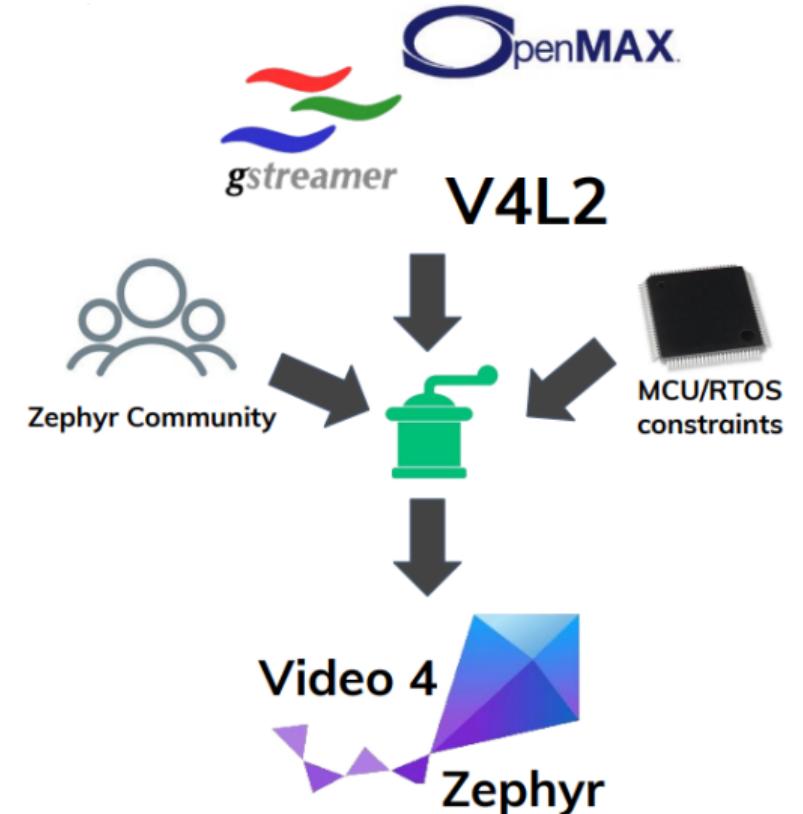
OpenCV → ???

PyTorch → ???

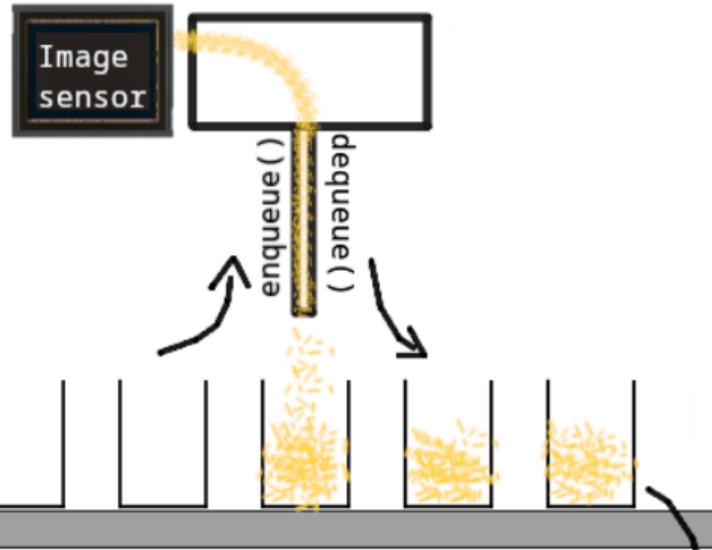
/dev/video0 → ???

These software assumes a lot of resources. Everything to reinvent!

Zephyr Video APIs

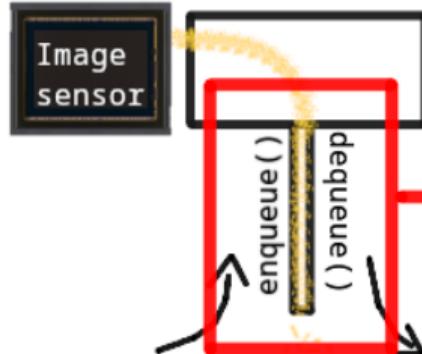


MIPI driver



Recycling buffers:
not reallocating: reusing

MIPI driver



Data API

not implemented
by all drivers

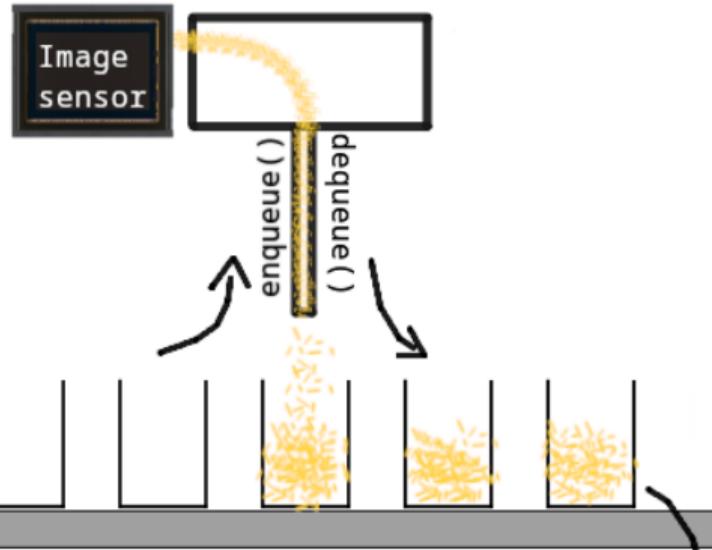


Recycling buffers:
not reallocating: reusing

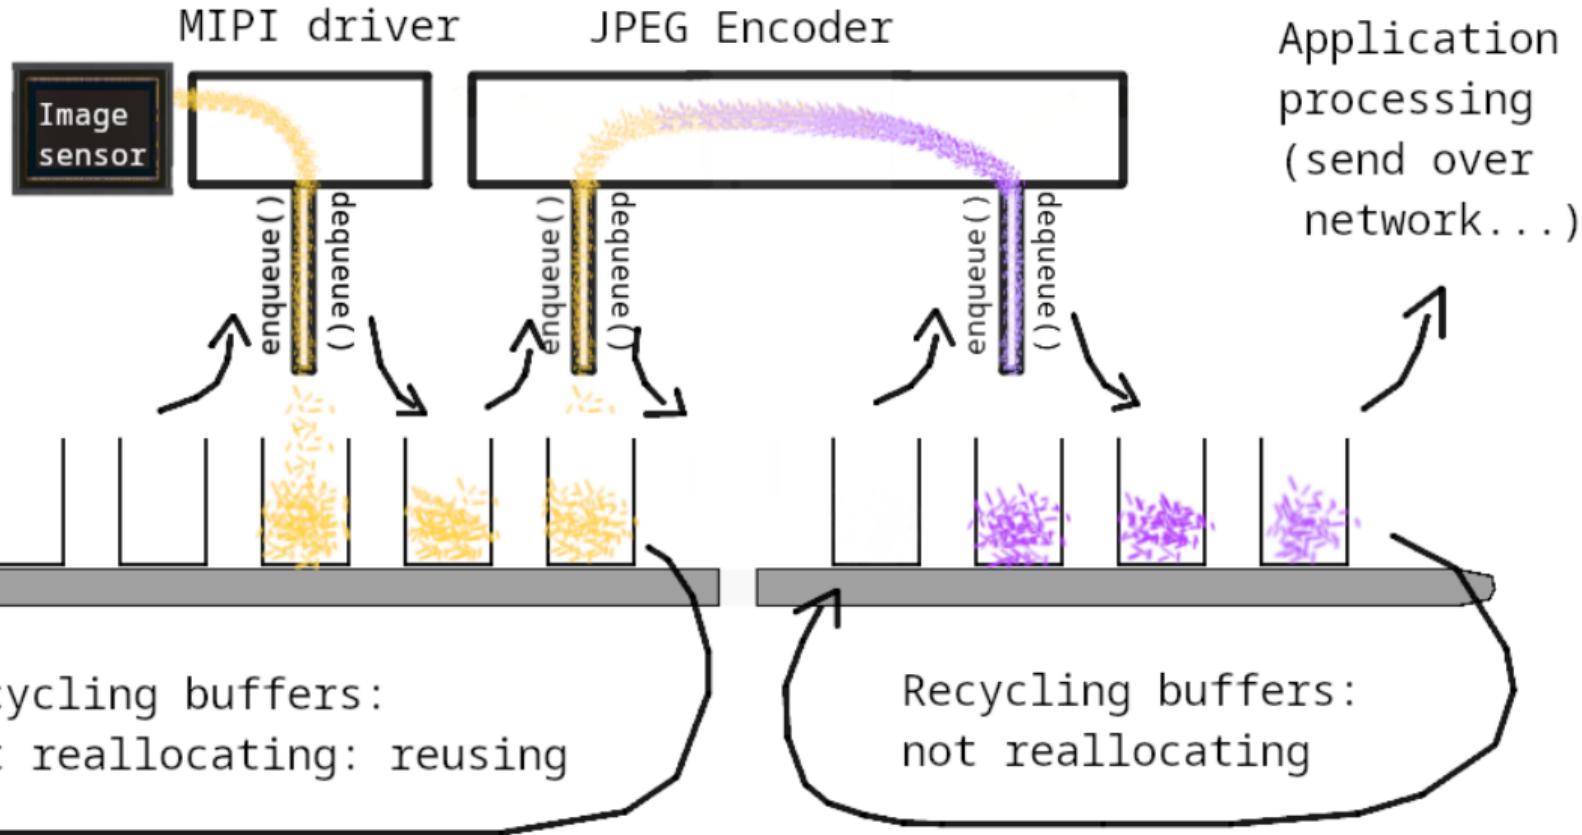
```
/* Error handling omitted */

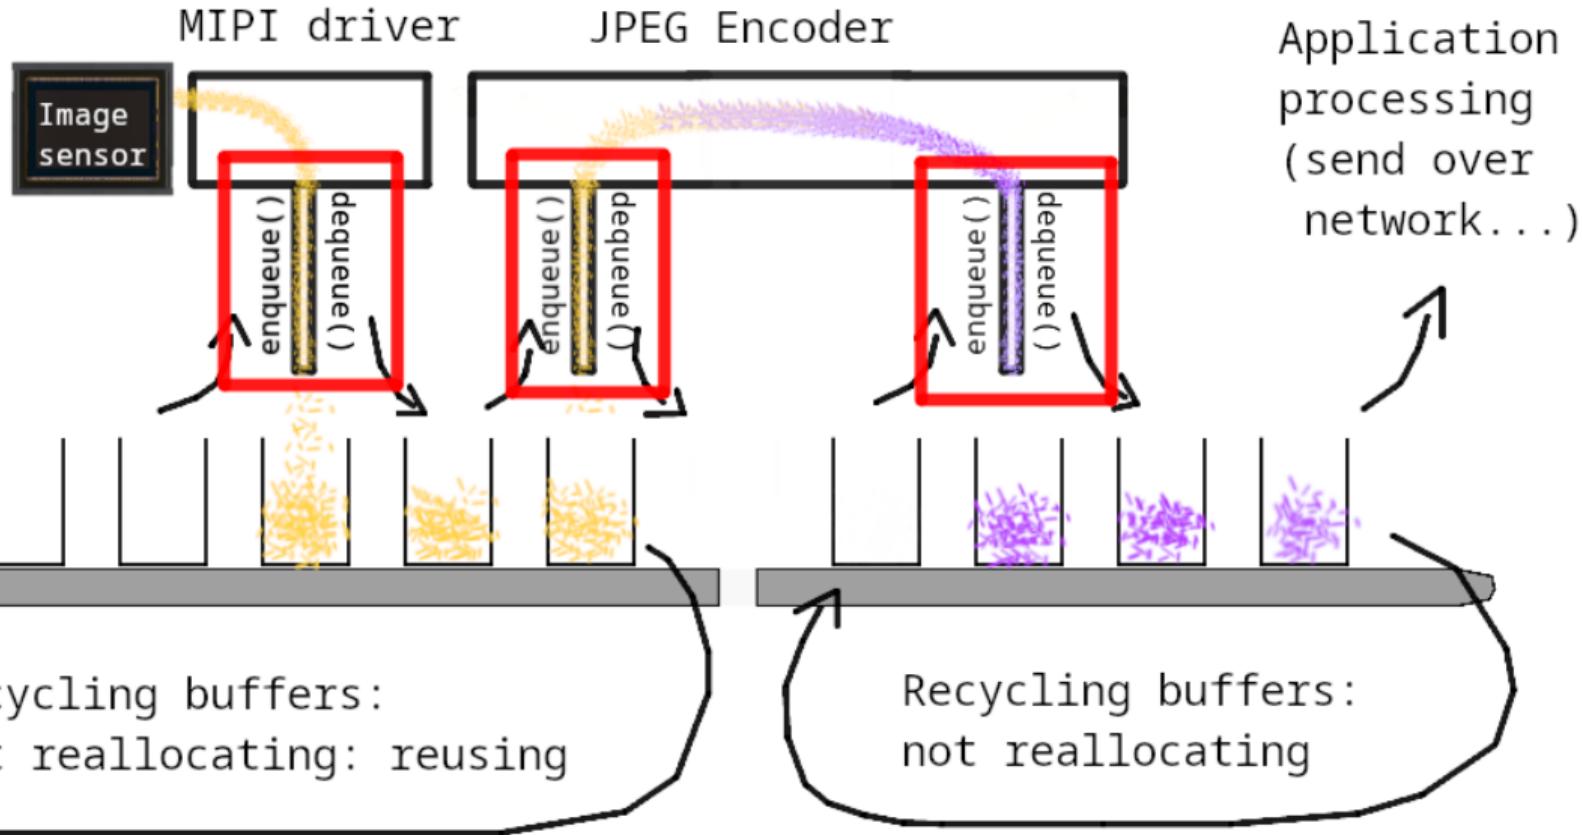
video_stream_start(mipi_dev);
vbuf = video_buffer_alloc(1280 * 720 * 3, K_FOREVER);
while (true) {
    video_enqueue(mipi_dev, VIDEO_EP_OUT, vbuf);
    video_dequeue(mipi_dev, VIDEO_EP_OUT, &vbuf, K_FOREVER);
    /* do something with vbuf->data */
}
```

MIPI driver



Recycling buffers:
not reallocating: reusing





```
/* Only one buffer of each */
vbuf_raw = video_buffer_alloc(1280 * 720 * 3, K_FOREVER);
vbuf_jpeg = video_buffer_alloc(JPEG_BUF_SIZE, K_FOREVER);
while (true) {
    video_enqueue(mipi_dev, VIDEO_EP_OUT, vbuf_raw);
    video_dequeue(mipi_dev, VIDEO_EP_OUT, &vbuf_raw, K_FOREVER);
    video_enqueue(jpeg_dev, VIDEO_EP_IN, vbuf_raw);
    video_dequeue(jpeg_dev, VIDEO_EP_IN, &vbuf_raw, K_FOREVER);
    video_enqueue(jpeg_dev, VIDEO_EP_OUT, vbuf_jpeg);
    video_dequeue(jpeg_dev, VIDEO_EP_OUT, &vbuf_jpeg, K_FOREVER);
    /* Do something with the JPEG data */
}
```

Nice and simple, but no concurrency: a lot of time spent waiting.

```
for (int i = 0; i < 3; i++) {
    vbuf_raw = video_buffer_alloc(1280 * 720 * 3, K_FOREVER);
    video_enqueue(mipi_dev, VIDEO_EP_OUT, vbuf_raw);
}

for (int i = 0; i < 3; i++) {
    vbuf_jpeg = video_buffer_alloc(JPEG_BUF_SIZE, K_FOREVER);
    video_enqueue(jpeg_dev, VIDEO_EP_OUT, vbuf_jpeg);
}

/* Thread 1: Handle the filled buffers */
while (true) {
    video_dequeue(mipi_dev, VIDEO_EP_OUT, &vbuf_raw, K_FOREVER);
    video_enqueue(jpeg_dev, VIDEO_EP_IN, vbuf_raw);
}

/* Thread 2: Handle the emptied buffers */
```

```
while (true) {
    video_dequeue(jpeg_dev, VIDEO_EP_IN, &vbuf_raw, K_FOREVER);
    video_enqueue(mipi_dev, VIDEO_EP_OUT, vbuf_raw);
}

/* Thread 3: Handle the JPEG buffers */
while (true) {
    video_dequeue(jpeg_dev, VIDEO_EP_IN, &vbuf_jpeg, K_FOREVER);
    app_use_jpeg_buffer(vbuf_jpeg);
    video_enqueue(jpeg_dev, VIDEO_EP_OUT, &vbuf_jpeg, K_FOREVER);
}
```

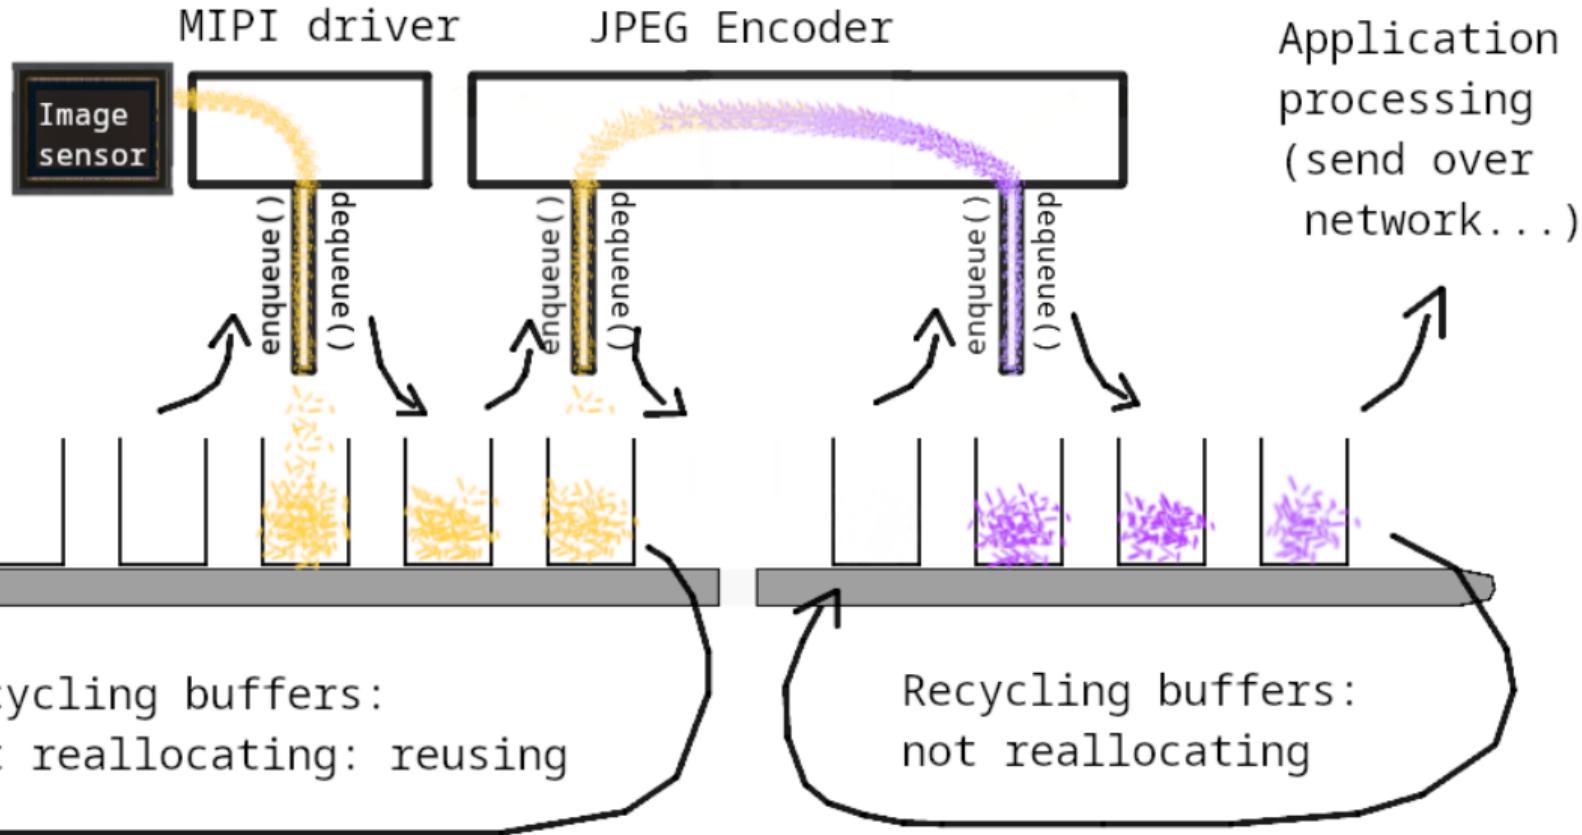
Fully parallel, but consumes a lot of threads: more memory overhead and context switching.

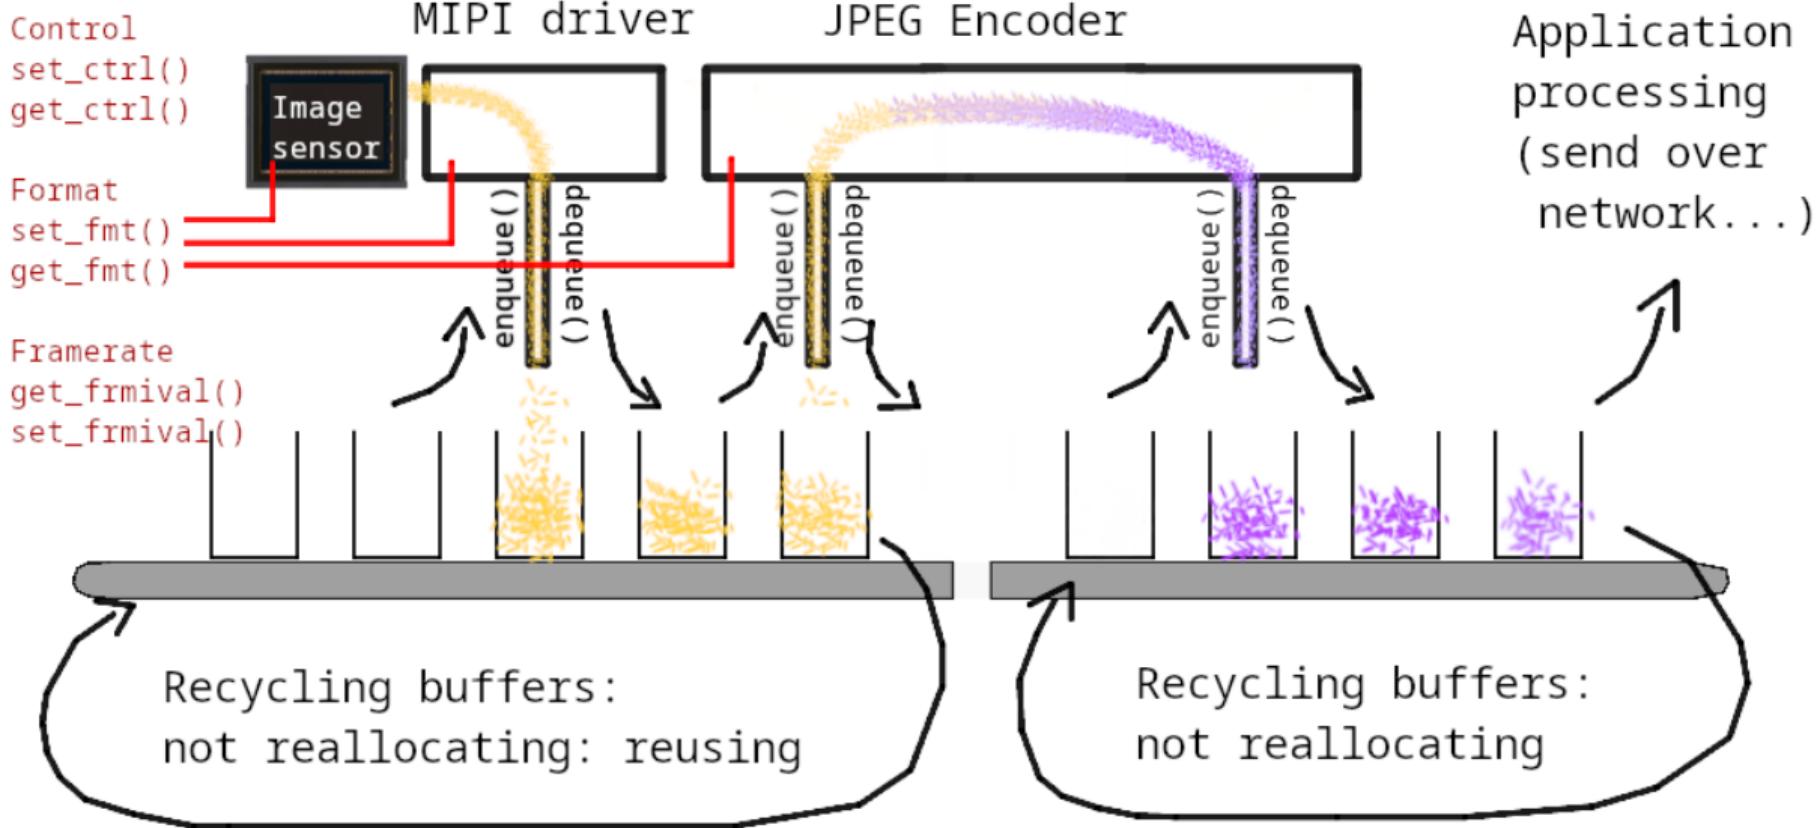
```
/* First prepare a few buffers */

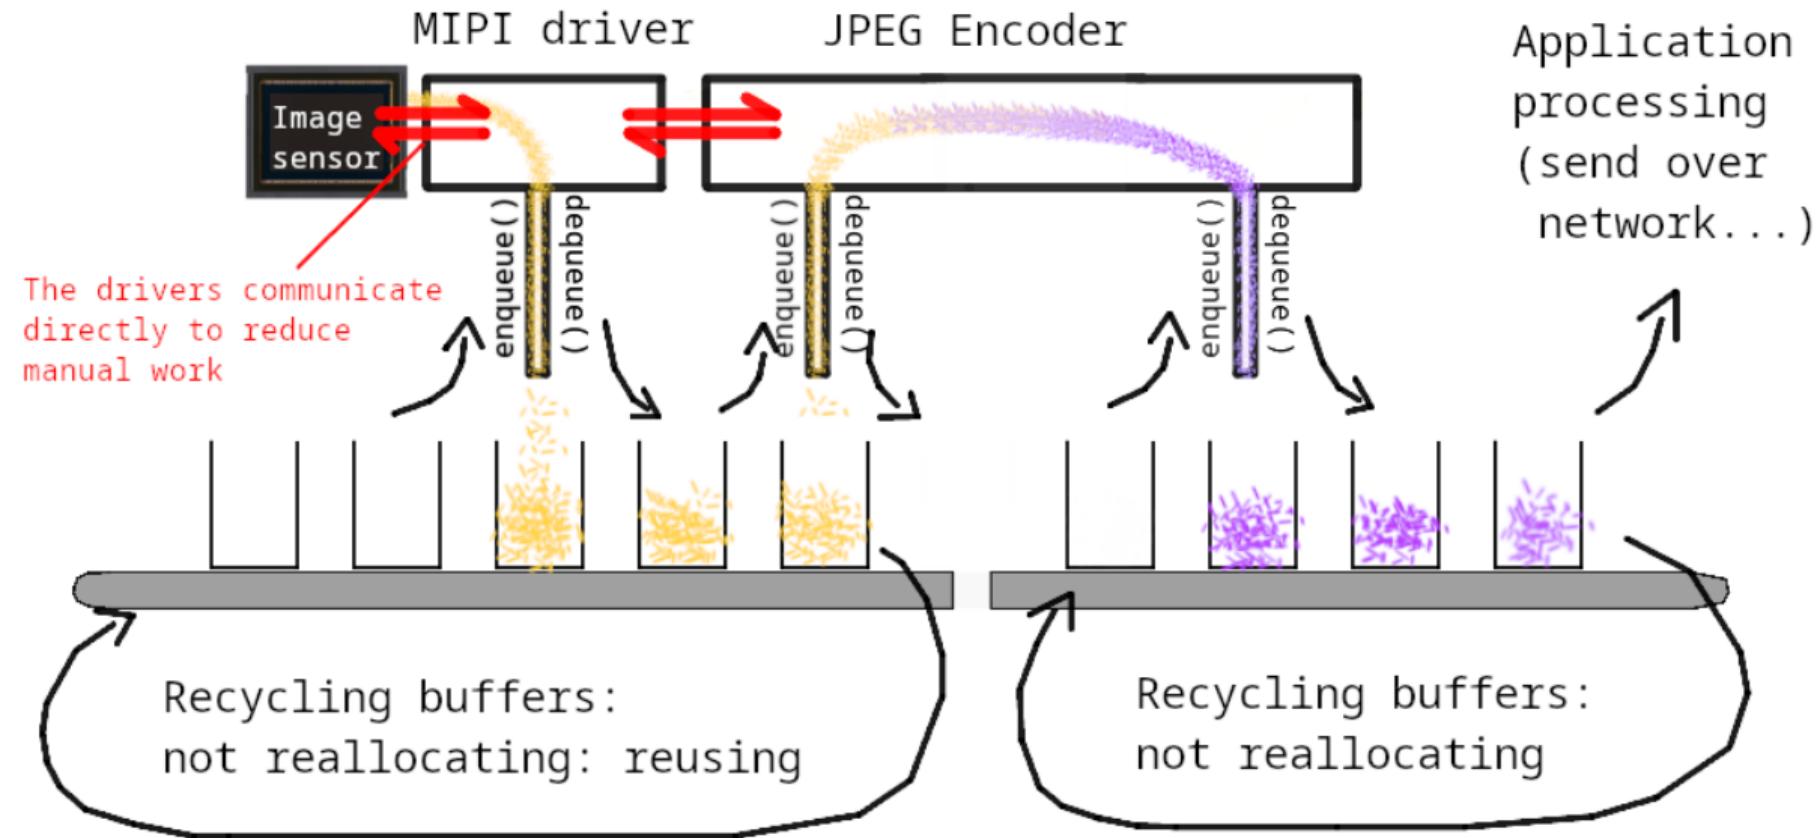
/* Subscribe to the events */
video_set_signal(mipi_dev, &signal);
video_set_signal(jpeg_dev, &signal);

/* React to the events */
while (k_poll(&events, 1, K_FOREVER) == 0) {
    if (video_dequeue(mipi_dev, VIDEO_EP_OUT, &vbuf_raw, K_NO_WAIT) == 0) {
        video_enqueue(jpeg_dev, VIDEO_EP_IN, vbuf_raw);
    }
    if (video_dequeue(jpeg_dev, VIDEO_EP_IN, &vbuf_raw, K_NO_WAIT) == 0) {
        video_enqueue(mipi_dev, VIDEO_EP_OUT, vbuf_raw);
    }
    if (video_dequeue(jpeg_dev, VIDEO_EP_OUT, &vbuf_jpeg, K_NO_WAIT) == 0) {
        video_enqueue(jpeg_dev, VIDEO_EP_OUT, vbuf_jpeg);
        app_use_jpeg_buffer();
    }
}
```

Good parallelism, single thread (more scalable). Maybe this can be automated for easier pipelines.







```
imx219: imx219@10 {
    compatible = "sony,imx219";
    port {
        imx219_ep_out: endpoint {
            remote-endpoint-label = "mipi0_ep_in";
        };
    };
};
```

```
mipi0: mipi@b1000010 {
    compatible = "tinyvision,mipi";
    port {
        mipi0_ep_in: endpoint {
            remote-endpoint-label = "imx219_ep_out";
        };
};
```

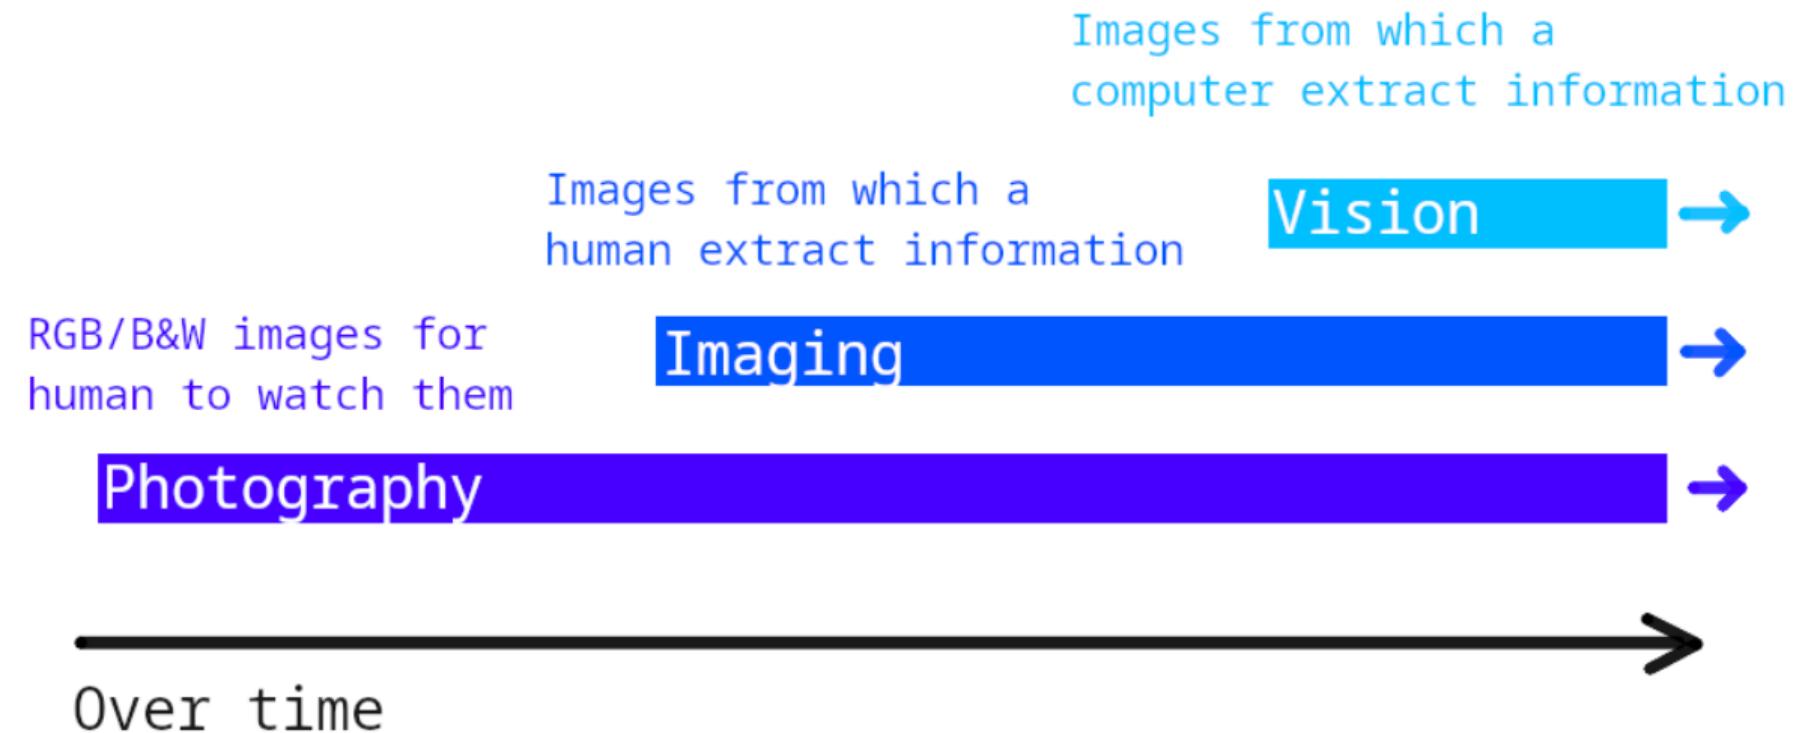
```
        mipi0_ep_out: endpoint {
            remote-endpoint-label = "jpegenc_ep_in";
        };
    };
};
```

```
jpegenc0: jpegenc@b1000010 {
    compatible = "tinyvision,jpegenc";
    port {
        jpegenc0_ep_in: endpoint {
            remote-endpoint-label = "mipi0_ep_in";
        };
        /* jpegenc0_ep_out: to the application */
    };
};
```

Next steps: automating more of the pipeline? What would make it more convenient?

Systems doing what?

On a journey from Phontons to Video, and how that is used



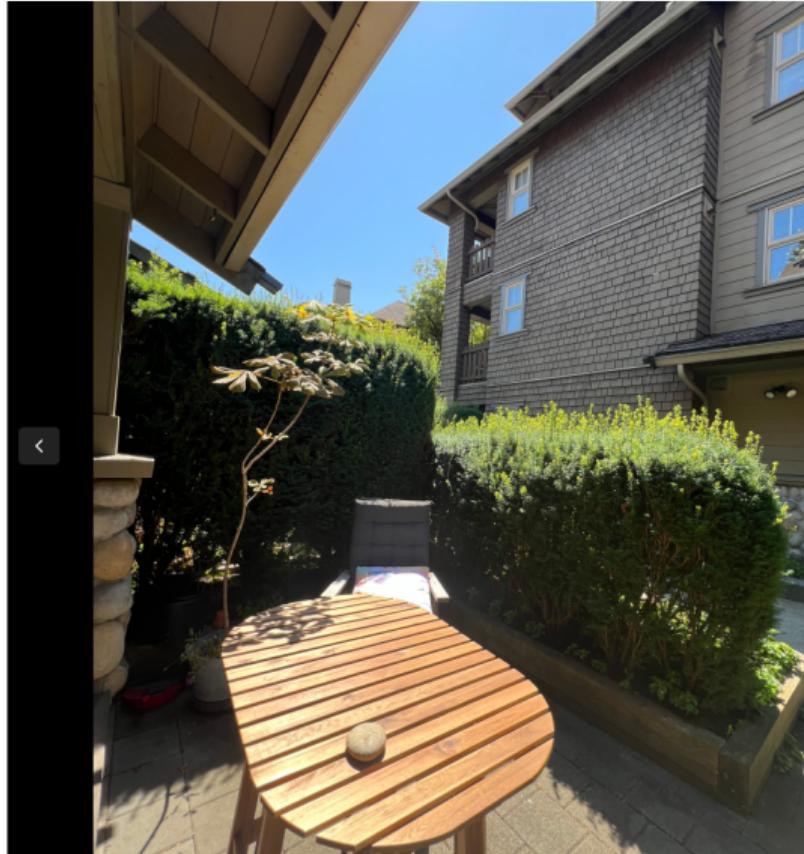
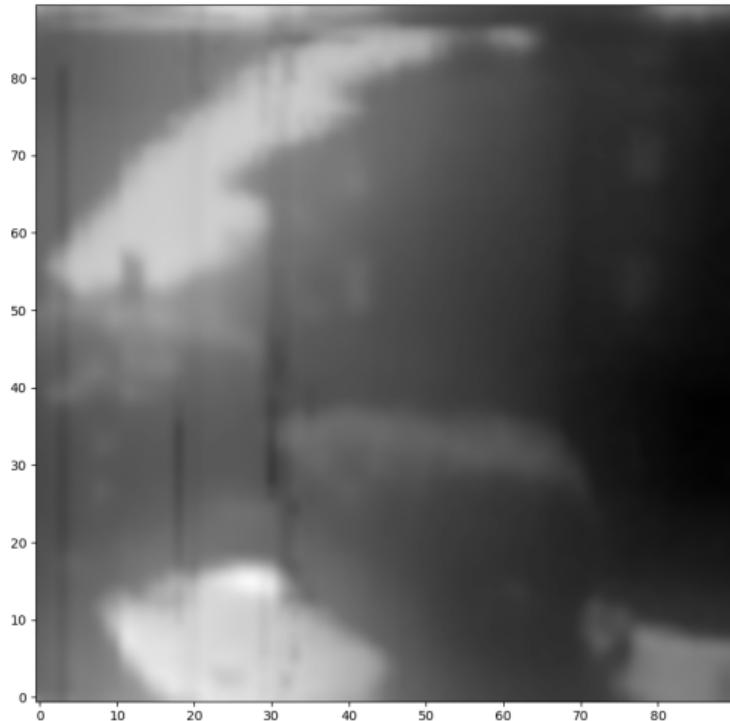
Inspired from a MIPI alliance presentation (2018)

Photodiode

Phenomenon of semiconductors producing voltage when exposed to the light.







Note: photoresistor instead of photodiode here

```
#include <zephyr/drivers/pwm.h> // if using servomotors
#include <zephyr/drivers/stepper.h> // if using stepper motors
#include <zephyr/drivers/adc.h> // measure the light intensity
```

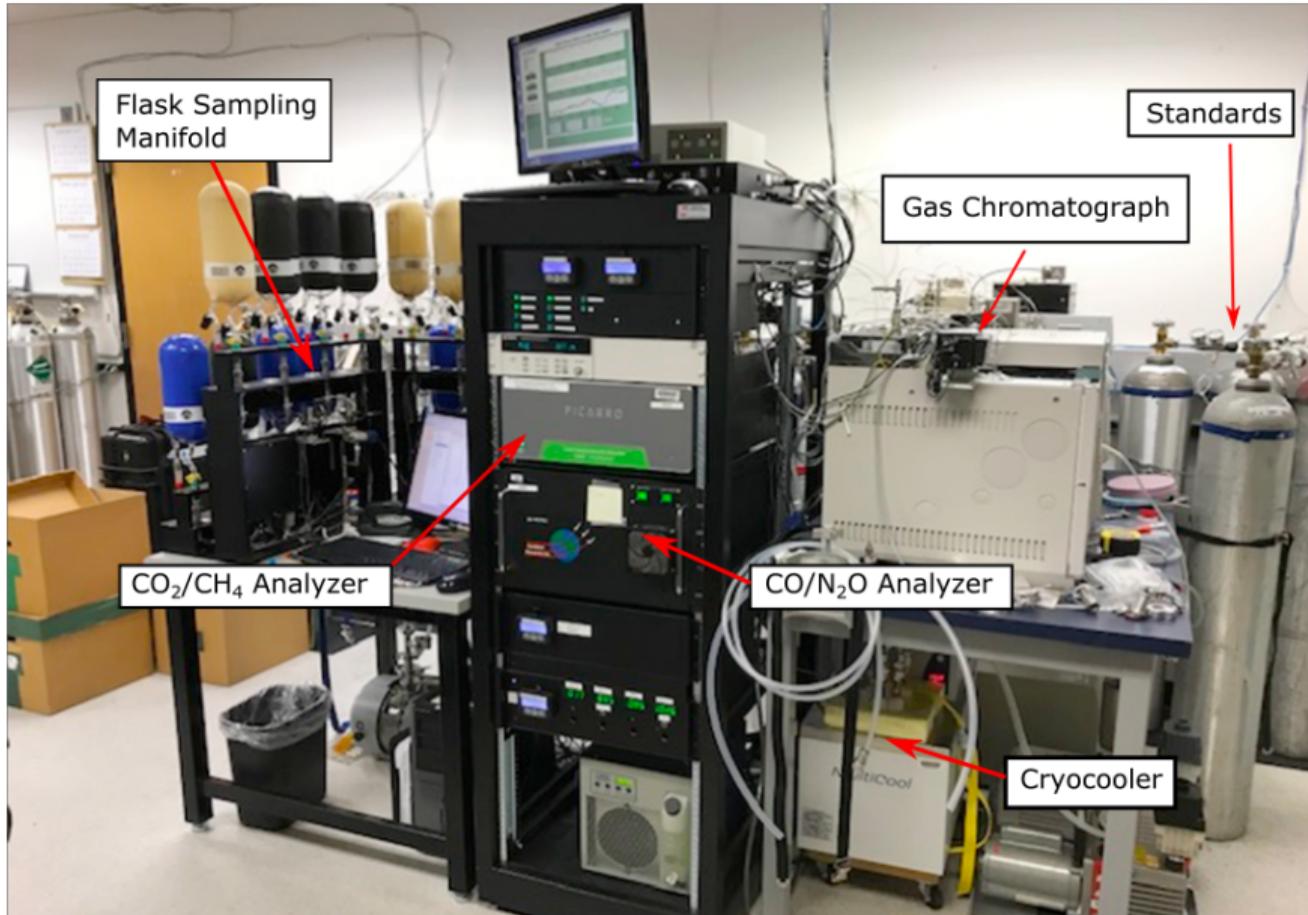
Photons → Photonics

Much more than just video:

→ Gas detection/characterization, i.e. NDIR CO₂ sensors

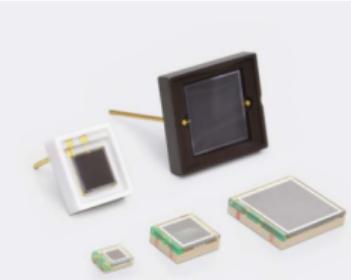
Industrial, safety, medical use-cases.

Since 1958: measuring Earth atmospheric CO₂ with "1-pixel image sensors"



→ Biology/medical research, i.e. DNA sequencing

MPPC® (multi-pixel photon counter)



S13360 series

MPPCs for precision measurement

MPPC is a type of device called SiPM (silicon photomultipliers). It is a new type of photon counting device that consists of multiple Geiger mode APD (avalanche photodiode) pixels. It is an opto-semiconductor with outstanding photon counting capability and low operating voltage and is immune to the effects of magnetic fields.

The S13360 series are MPPCs for precision measurement. The MPPCs inherit the superb low afterpulse characteristics of previous products and further provide lower crosstalk and lower dark count. They are suitable for precision measurement, such as flow cytometry, DNA sequencer, laser microscope, and fluorescence measurement, that requires low noise characteristics.

Features

- Reduced crosstalk and dark count
(compared to previous products)
- Outstanding photon counting capability (outstanding photon detection efficiency versus numbers of incident photons)
- Compact
- Operates at room temperature
- Low voltage ($V_{DD} = 52 \text{ V}$ typ.) operation

Applications

- Fluorescence measurement
- Laser microscopes
- Flow cytometry
- **DNA sequencers**
- Environmental analysis
- Various academic research

Sensing voltage: not a very Linux thing to do...

A single line of pixels

Line sensors: a single line at a time.

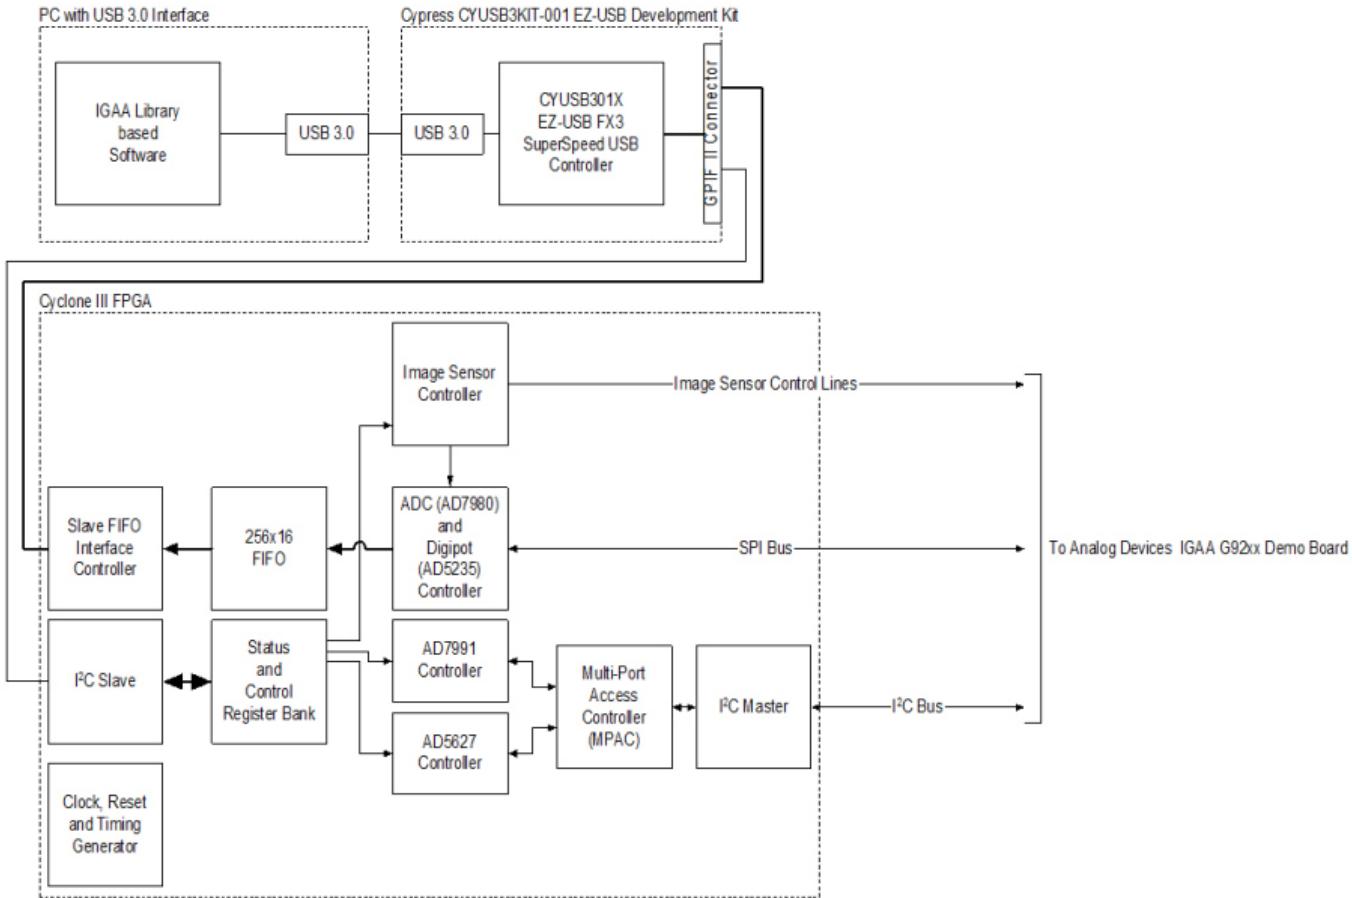
Not digital interface but analog interface: need an ADC to get digital readout.

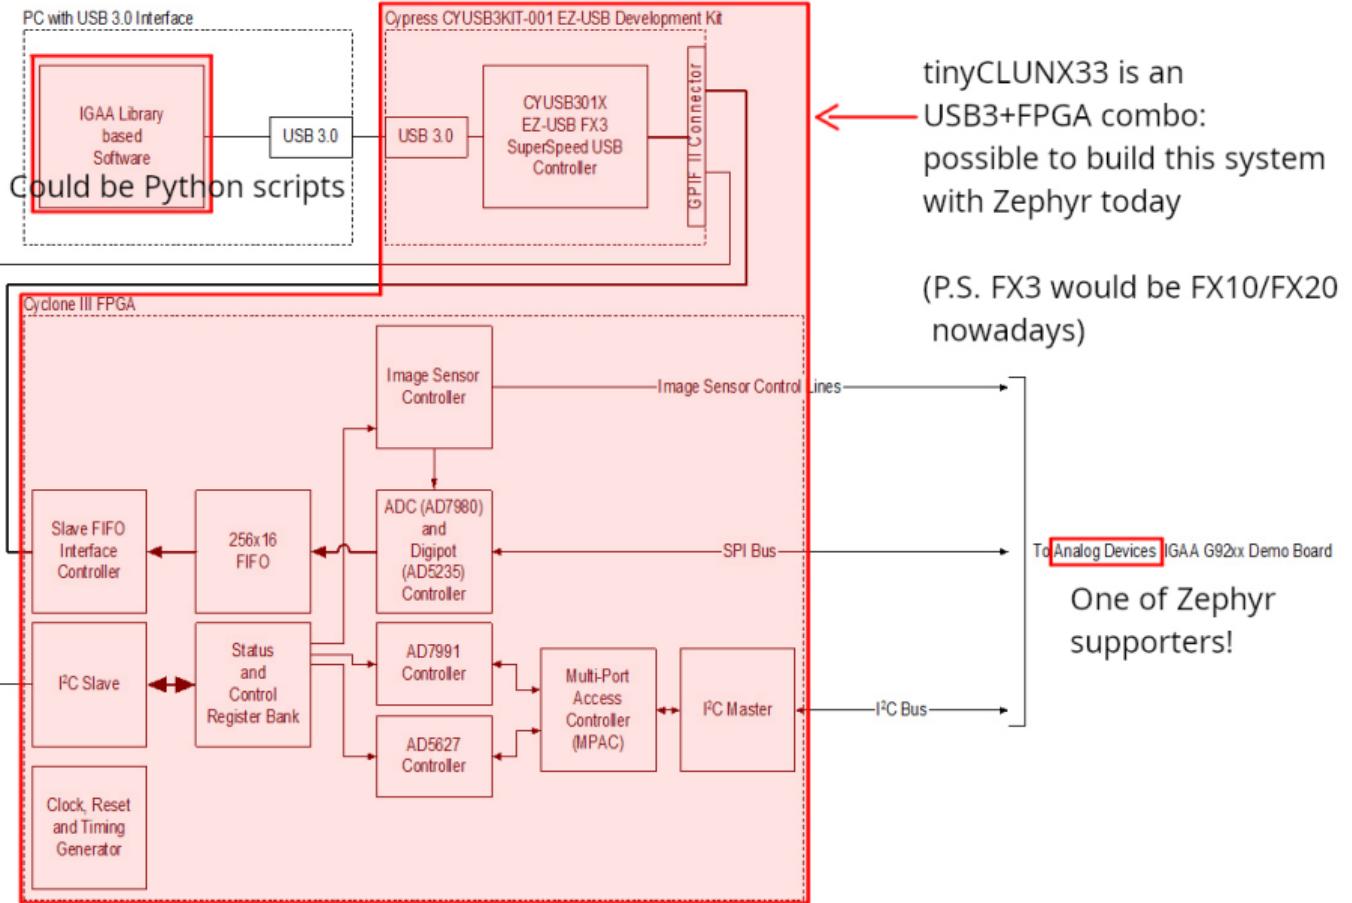
Sensing one pixel at a time, need to go fast!



Requires a fast ADC to handle the input.

Hamamatsu recommends an FPGA (customizable chip) + Analog Devices Analog-Digital-Converter (ADC).





Multiple lines

Tools that can be used for building video systems: hardware to access the sensors implement all of that chain

- Difficulty of embedded video: accessing parallel port or MIPI
- Can use adapter chips like Himax HX6538 (not yet supported) or small FPGAs

What comes out of an image sensor

Dark (no auto-exposure) Green (no color correction)

Steps of an ISP.

Why an ISP is useful for robotics?

- Get always values within same range
- Poor exposure: no data at all
- De-fisheye
- Avoid artefacts to trigger a detection on the NPU or other vision algorithm

Conclusion: A lot to handle to get a reasonable image out of a sensor!

Hardware that can help accessing this image.

What it takes...

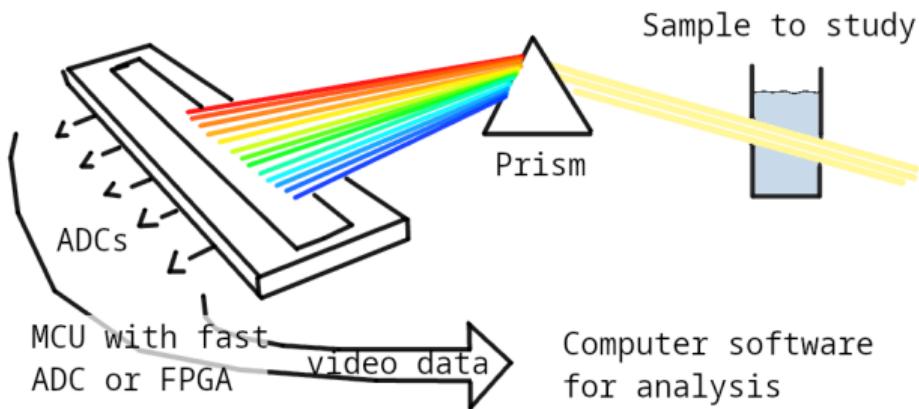
What would it take to build various devices on Zephyr

!! disclaimer: hardware is hard !!

!! disclaimer: not everything shown has drivers !!

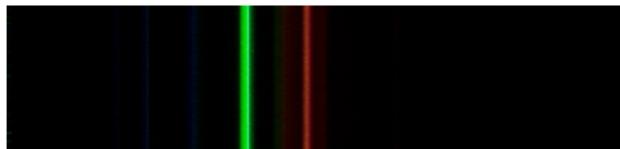
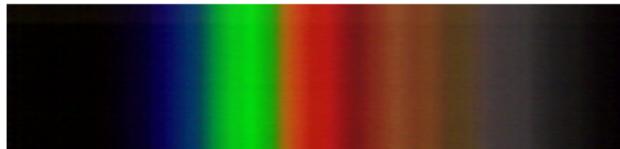
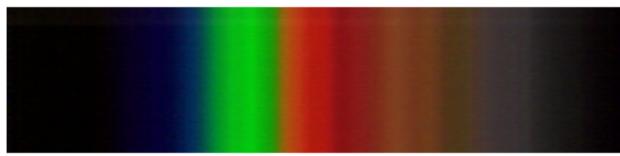
Also works without Zephyr, just putting things in perspective.

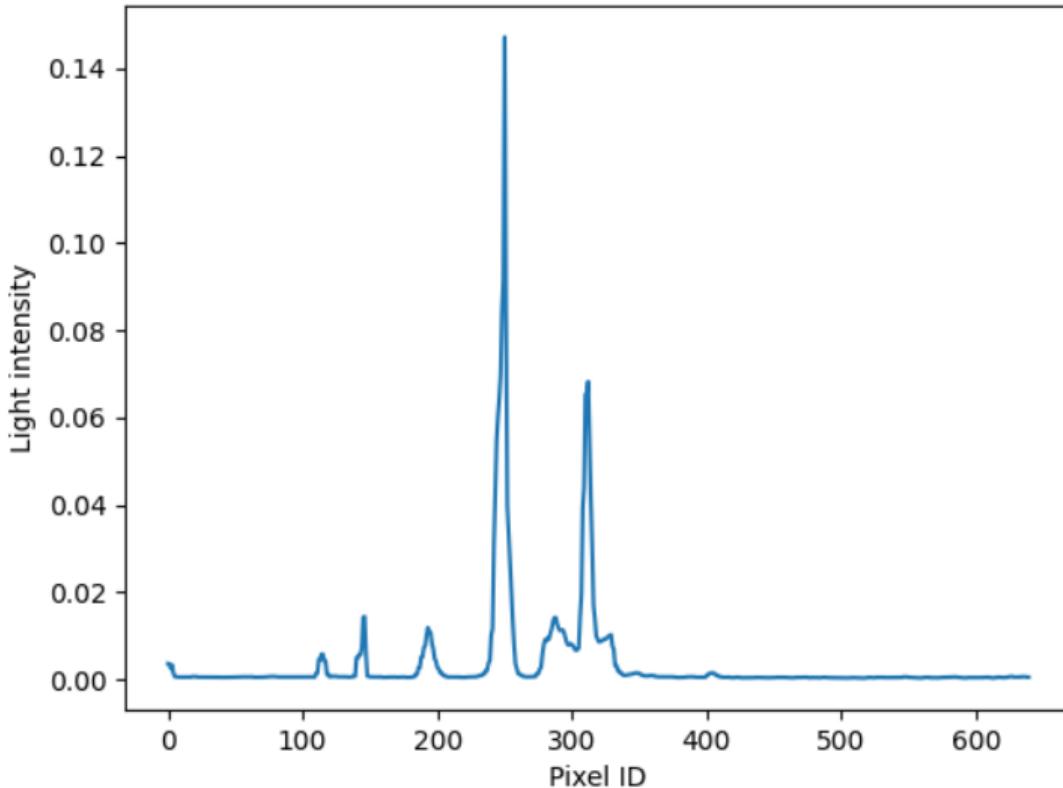
What it takes... Spectrophotometer



Need a very fast ADC! Not many board will have one...

→ Good to have a lot of options.





What it takes... Drones

VMU RT1170

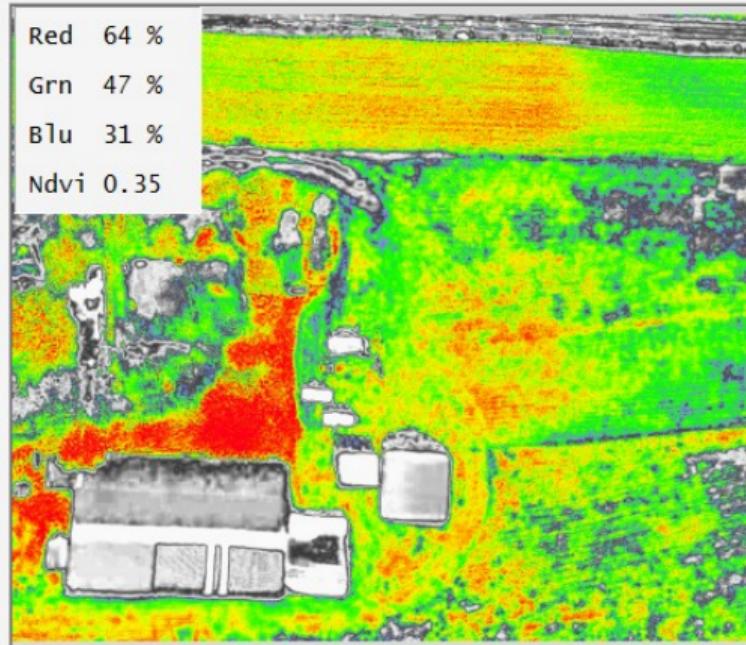
Overview

The VMU RT1170 features an i.MX RT1176 dual core MCU with the Cortex-M7 core at 1 GHz and a Cortex-M4 at 400 MHz. The i.MX RT1176 MCU offers support over a wide temperature range and is qualified for consumer, industrial and automotive markets. The VMU RT1170 is the default VMU for CogniPilot's Cerebri, a Zephyr RTOS based Autopilot.

Hardware

- MIMXRT1176DVMAA MCU
 - 1GHz Cortex-M7 & 400Mhz Cortex-M4
 - 2MB SRAM with 512KB of TCM for Cortex-M7 and 256KB of TCM for Cortex-M4
- Memory
 - 512 Mbit Octal Flash
 - TF socket for SD card
- Ethernet
 - 2 wire 100BASE-T1
- USB
 - USB 2.0 connector
- Power
 - Redundant dual picoflex power ports
- Debug
 - 10 pin debug and shell adapter board to 20 Pin JTAG debugger and USB-C shell
- Sensors





Double click on images to open them with the default application

What it takes... Yeast monitoring station

Monitoring process of beer, kombucha, lactic fermentation

Video but also...

CO2 polling for building charts.

LED API for illuminating when taking a capture.

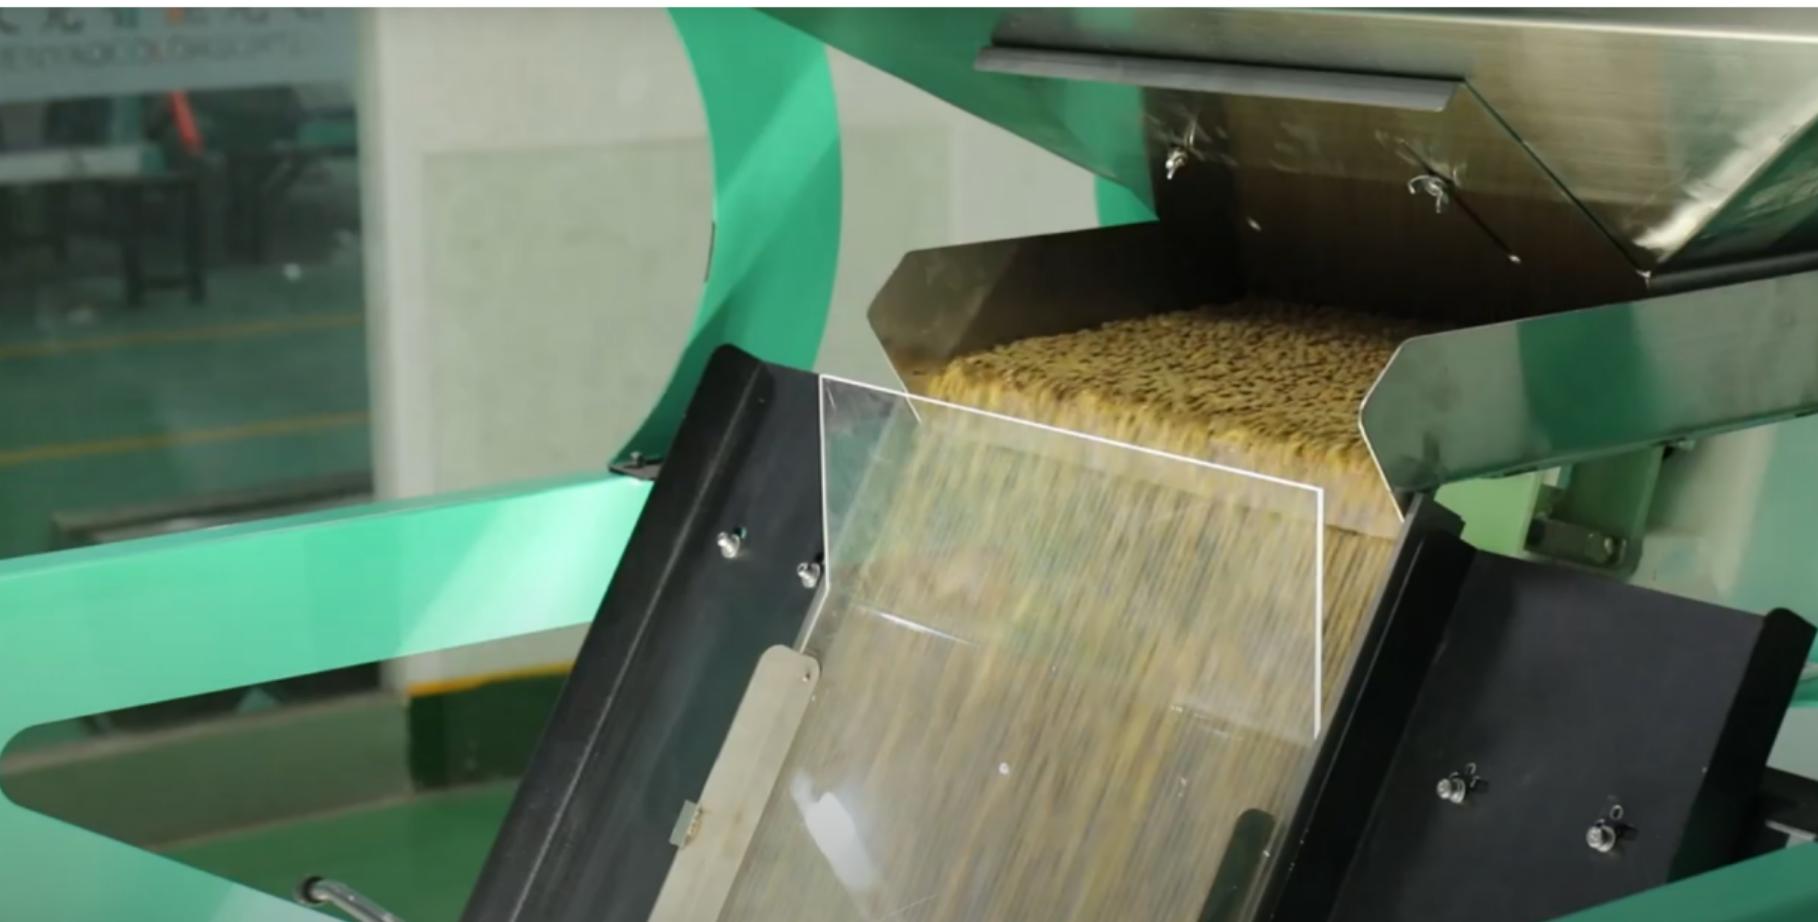
Video controls API for manual exposure tuning during "flash".

Wi-Fi to the home router.

HTTP client for sending the results.

What it takes... Sorting machine

Fastest board you can get! Most real-time you can get!



Elimination of the rotten and unripe berries and the stems via pneumatic injectors.



Blowing air on everything not looking like a fresh grain of raisin.

16x slow motion



What it takes... Endoscopes/Borescope

Cameras used by surgeons

Example of real endoscope camera module (CAMEMAKE):



Image sensor driver

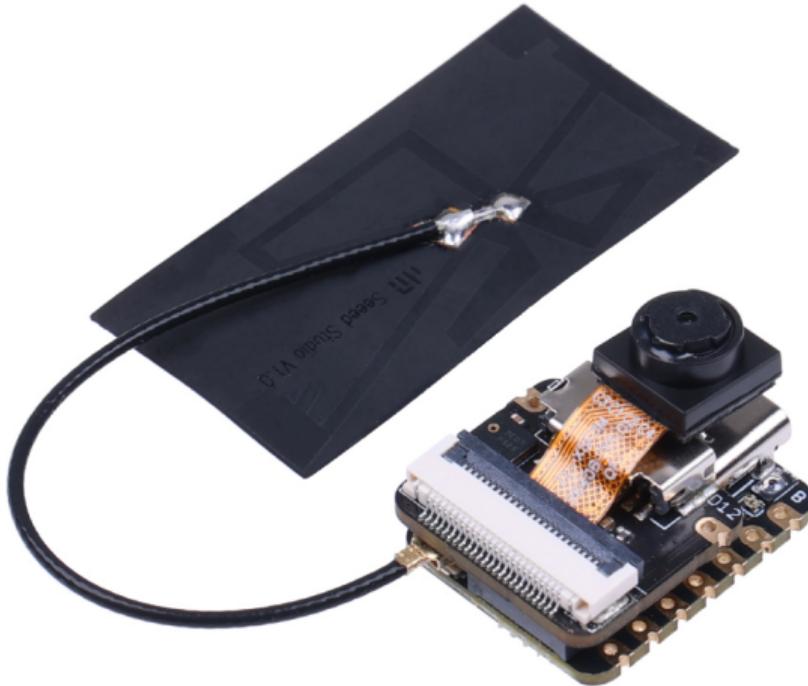
Video transfer driver <-> Video APIs <-> UVC

Parallel
port



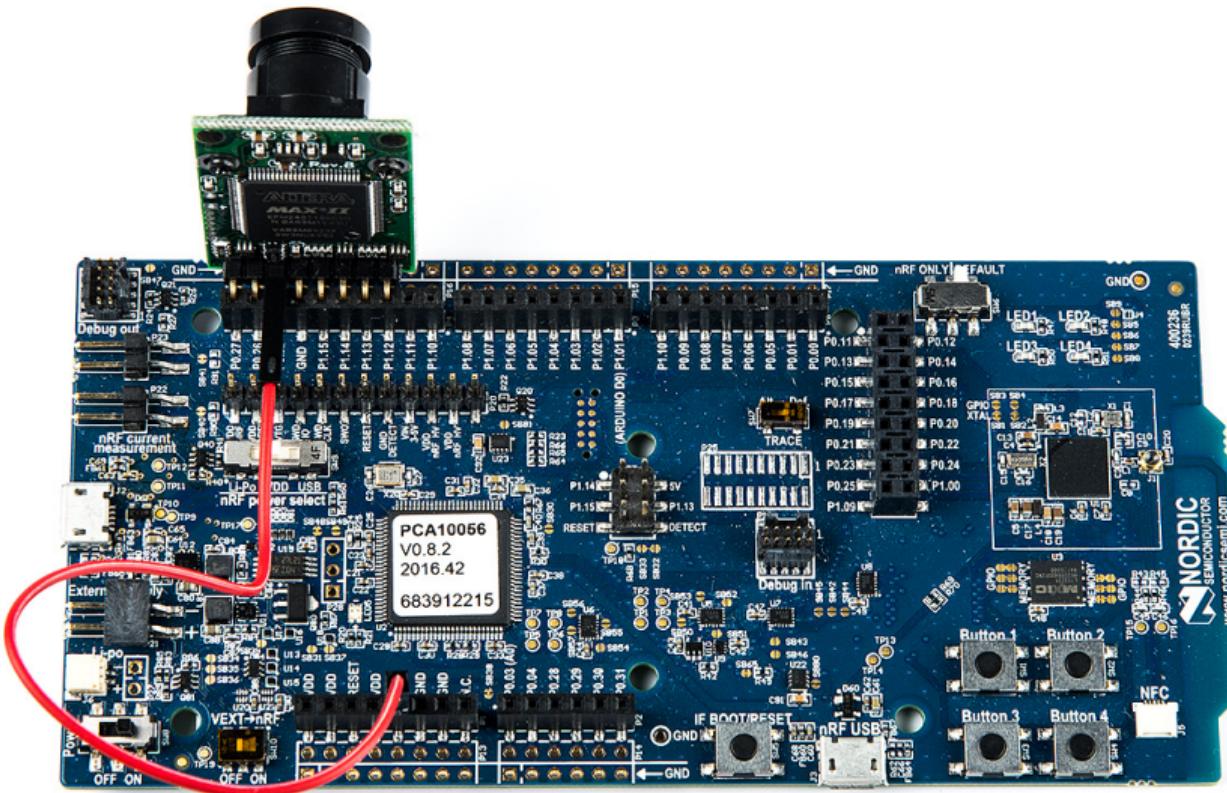
Image
sensor

What it takes... Wi-Fi Smartglasses



What it takes... Bluetooth Smartglasses

Pre-Zephyr Nordic era: needs conversion.



Any MCU with
Bluetooth

↔ SPI →

Himax HM0360
(not yet supported)



640x480

Beyond Zephyr: ecosystem around it

What UVC adds to the table:

- USB camera protocol supported on Linux, Windows, MaxOS, Android, iPad (not iOS yet), BSDs, 9front, QNX... (thanks to laptop lid cameras)
- Linux interoperability: Standardize all the video controls with Linux
- ROS2: integration of robotics (via USB cameras)
- OpenCV (via USB cameras)
- Want to support a new sensor on any ecosystem? Bring Zephyr support, and now it's everywhere